

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, cross_val_score

from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.svm import SVR

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
df = pd.read_csv("cars.csv")
df.head()
```

Out[2]:

	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	engine- type
0	3	?	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
1	3	?	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
2	1	?	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	ohcv
3	2	164	audi	gas	sedan	fwd	front	66.2	54.3	ohc
4	2	164	audi	gas	sedan	4wd	front	66.4	54.3	ohc

In [3]:

```
#Checking the number of nan values
```

In [4]:

```
df.isna().sum()
```

Out[4]:

```

symboling          0
normalized-losses  0
make              0
fuel-type         0
body-style        0
drive-wheels      0
engine-location   0
width            0
height           0
engine-type       0
engine-size       0
horsepower        0
city-mpg          0
highway-mpg       0
price            0
dtype: int64

```

In [5]:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling             205 non-null   int64
1   normalized-losses     205 non-null   object
2   make                 205 non-null   object
3   fuel-type            205 non-null   object
4   body-style           205 non-null   object
5   drive-wheels         205 non-null   object
6   engine-location       205 non-null   object
7   width                205 non-null   float64
8   height               205 non-null   float64
9   engine-type          205 non-null   object
10  engine-size           205 non-null   int64
11  horsepower            205 non-null   object
12  city-mpg              205 non-null   int64
13  highway-mpg           205 non-null   int64
14  price                205 non-null   int64
dtypes: float64(2), int64(5), object(8)
memory usage: 24.1+ KB

```

In [6]:

```
df["normalized-losses"].value_counts()
```

Out[6]:

?	41
161	11
91	8
150	7
104	6
128	6
134	6
85	5
94	5
65	5
168	5
102	5
74	5
103	5
95	5
106	4
118	4
148	4
93	4
122	4
115	3
125	3
154	3
101	3
83	3
137	3
192	2
129	2
89	2
153	2
81	2
87	2
197	2
194	2
158	2
108	2
145	2
164	2
113	2
119	2
110	2
188	2
256	1
107	1
98	1
142	1
78	1
77	1
90	1
121	1
231	1
186	1

Name: normalized-losses, dtype: int64

In [7]:

```
df["horsepower"].value_counts()
```

Out[7]:

68	19
70	11
69	10
116	9
110	8
95	7
88	6
160	6
101	6
114	6
62	6
82	5
76	5
97	5
102	5
84	5
145	5
111	4
86	4
92	4
123	4
182	3
121	3
85	3
207	3
152	3
73	3
90	3
176	2
52	2
100	2
184	2
56	2
?	2
162	2
94	2
161	2
112	2
156	2
155	2
72	1
48	1
200	1
106	1
60	1
78	1
115	1
288	1
120	1
142	1
143	1
140	1
175	1

```
134    1
55     1
135    1
64     1
262    1
154    1
58     1
```

Name: horsepower. dtype: int64

## Handling the Nan values

In [8]:

```
#Replacing the values with np.nan
df["normalized-losses"].replace("?", np.nan, inplace=True)
df["horsepower"].replace("?", np.nan, inplace=True)

#Changing the datatype
df["normalized-losses"] = df["normalized-losses"].astype("float")
df["horsepower"] = df["horsepower"].astype("float")

#Getting the mean value
nmean = df["normalized-losses"].mean()
hmean = df["horsepower"].mean()

#Filling the missing values with mean values
df["normalized-losses"].fillna(nmean, inplace=True)
df["horsepower"].fillna(hmean, inplace=True)
```

In [9]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null   int64
1   normalized-losses      205 non-null   float64
2   make                   205 non-null   object
3   fuel-type              205 non-null   object
4   body-style             205 non-null   object
5   drive-wheels           205 non-null   object
6   engine-location        205 non-null   object
7   width                  205 non-null   float64
8   height                 205 non-null   float64
9   engine-type            205 non-null   object
10  engine-size            205 non-null   int64
11  horsepower             205 non-null   float64
12  city-mpg               205 non-null   int64
13  highway-mpg            205 non-null   int64
14  price                  205 non-null   int64
dtypes: float64(4), int64(5), object(6)
memory usage: 24.1+ KB
```


- It is observed that the datatype of columns containing numeric values now have proper datatype.
- Also the operation of handling nan values is performed successfully.

# Outliers

- Extracting the outliers with the help of boxplot.

In [10]:

```
plt.figure(figsize=(12,8))
sns.boxplot(data=df , x="price" , y="make")
plt.grid(True)
plt.show()
```



## Dropping the rows containing outliers and cleaning the data.

In [11]:

```
df[(df["make"]=="dodge") & (df["price"]>11000)]
```

Out[11]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
29	3	145.0	dodge	gas	hatchback	fwd	front	66.3	50.2	ohc

In [12]:

```
df.drop(29, inplace=True)
```

In [13]:

```
df[(df["make"]=="honda") & (df["price"]>11000)]
```

Out[13]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
41	0	85.0	honda	gas	sedan	fwd	front	65.2	54.1	ohc	

In [14]:

```
df.drop(41, inplace=True)
```

In [15]:

```
df[(df["make"]=="isuzu") & (df["price"]>15000)]
```

Out[15]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
45	0	122.0	isuzu	gas	sedan	fwd	front	63.6	52.0	ohc	

In [16]:

```
df.drop(45, inplace=True)
```

In [17]:

```
df[(df["make"]=="mitsubishi") & (df["price"]>13000)]
```

Out[17]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
83	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	oh	
84	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	oh	

In [18]:

```
df.drop([83,84], inplace=True)
```

In [19]:

```
df[(df["make"]=="plymouth") & (df["price"]>10000)]
```

Out[19]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
124	3	122.0	plymouth	gas	hatchback	rwd	front	66.3	50.2	oh

In [20]:

```
df.drop(124, inplace=True)
```

In [21]:

```
df[(df["make"]=="toyota") & (df["price"]>12000)]
```

Out[21]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
172	2	134.0	toyota	gas	convertible	rwd	front	65.6	53.0	ohc
178	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	dohc
179	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	dohc
180	-1	90.0	toyota	gas	sedan	rwd	front	66.5	54.1	dohc
181	-1	122.0	toyota	gas	wagon	rwd	front	66.5	54.1	dohc

In [22]:

```
df.drop([172,178,179,180,181], inplace=True)
```

## Encoding the cleaned data.

In [23]:

```
df_cat = df.select_dtypes(object)
df_num = df.select_dtypes(["int64", "float64"])
```



In [24]:

df\_num

Out[24]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg	price
0	3	122.0	64.1	48.8	130	111.0	21	27	13495
1	3	122.0	64.1	48.8	130	111.0	21	27	16500
2	1	122.0	65.5	52.4	152	154.0	19	26	16500
3	2	164.0	66.2	54.3	109	102.0	24	30	13950
4	2	164.0	66.4	54.3	136	115.0	18	22	17450
...	...	...	...	...	...	...	...	...	...
200	-1	95.0	68.9	55.5	141	114.0	23	28	16845
201	-1	95.0	68.8	55.5	141	160.0	19	25	19045
202	-1	95.0	68.9	55.5	173	134.0	18	23	21485
203	-1	95.0	68.9	55.5	145	106.0	26	27	22470
204	-1	95.0	68.9	55.5	141	114.0	19	25	22625

194 rows × 9 columns

In [25]:

df\_cat

Out[25]:

	make	fuel-type	body-style	drive-wheels	engine-location	engine-type
0	alfa-romero	gas	convertible	rwd	front	dohc
1	alfa-romero	gas	convertible	rwd	front	dohc
2	alfa-romero	gas	hatchback	rwd	front	ohcv
3	audi	gas	sedan	fwd	front	ohc
4	audi	gas	sedan	4wd	front	ohc
...	...	...	...	...	...	...
200	volvo	gas	sedan	rwd	front	ohc
201	volvo	gas	sedan	rwd	front	ohc
202	volvo	gas	sedan	rwd	front	ohcv
203	volvo	diesel	sedan	rwd	front	ohc
204	volvo	gas	sedan	rwd	front	ohc

194 rows × 6 columns

In [26]:

```
from sklearn.preprocessing import LabelEncoder
```

In [27]:

```
for col in df_cat:  
    le = LabelEncoder()  
    df_cat[col] = le.fit_transform(df_cat[col])
```

In [28]:

df\_cat

Out[28]:

	make	fuel-type	body-style	drive-wheels	engine-location	engine-type
0	0	1	0	2	0	0
1	0	1	0	2	0	0
2	0	1	2	2	0	5
3	1	1	3	1	0	3
4	1	1	3	0	0	3
...	...	...	...	...	...	...
200	21	1	3	2	0	3
201	21	1	3	2	0	3
202	21	1	3	2	0	5
203	21	0	3	2	0	3
204	21	1	3	2	0	3

194 rows × 6 columns

- Successfully converted categorical data into numerical data using LabelEncoder.

## Creating new df by combining df\_cat & df\_num

In [29]:

df\_num

Out[29]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg	price
0	3	122.0	64.1	48.8	130	111.0	21	27	13495
1	3	122.0	64.1	48.8	130	111.0	21	27	16500
2	1	122.0	65.5	52.4	152	154.0	19	26	16500
3	2	164.0	66.2	54.3	109	102.0	24	30	13950
4	2	164.0	66.4	54.3	136	115.0	18	22	17450
...	...	...	...	...	...	...	...	...	...
200	-1	95.0	68.9	55.5	141	114.0	23	28	16845
201	-1	95.0	68.8	55.5	141	160.0	19	25	19045
202	-1	95.0	68.9	55.5	173	134.0	18	23	21485
203	-1	95.0	68.9	55.5	145	106.0	26	27	22470
204	-1	95.0	68.9	55.5	141	114.0	19	25	22625

194 rows × 9 columns

In [30]:

df\_cat

Out[30]:

	make	fuel-type	body-style	drive-wheels	engine-location	engine-type
0	0	1	0	2	0	0
1	0	1	0	2	0	0
2	0	1	2	2	0	5
3	1	1	3	1	0	3
4	1	1	3	0	0	3
...	...	...	...	...	...	...
200	21	1	3	2	0	3
201	21	1	3	2	0	3
202	21	1	3	2	0	5
203	21	0	3	2	0	3
204	21	1	3	2	0	3

194 rows × 6 columns

In [31]:

```
df = pd.concat([df_cat, df_num], axis=1)
```

In [32]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194 entries, 0 to 204
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   make                  194 non-null    int32
 1   fuel-type             194 non-null    int32
 2   body-style            194 non-null    int32
 3   drive-wheels          194 non-null    int32
 4   engine-location       194 non-null    int32
 5   engine-type           194 non-null    int32
 6   symboling             194 non-null    int64
 7   normalized-losses     194 non-null    float64
 8   width                 194 non-null    float64
 9   height                194 non-null    float64
10   engine-size           194 non-null    int64
11   horsepower            194 non-null    float64
12   city-mpg              194 non-null    int64
13   highway-mpg           194 non-null    int64
14   price                 194 non-null    int64
dtypes: float64(4), int32(6), int64(5)
memory usage: 19.7 KB
```

## EDA & Preprocessing


In [33]:

```
from scipy.stats import skew
```

In [34]:

```
for col in df:  
    print(col)  
    print(skew( df[col] ))  
  
    plt.figure()  
    sns.distplot(df[col])  
    plt.show()
```

0.21386866184357742



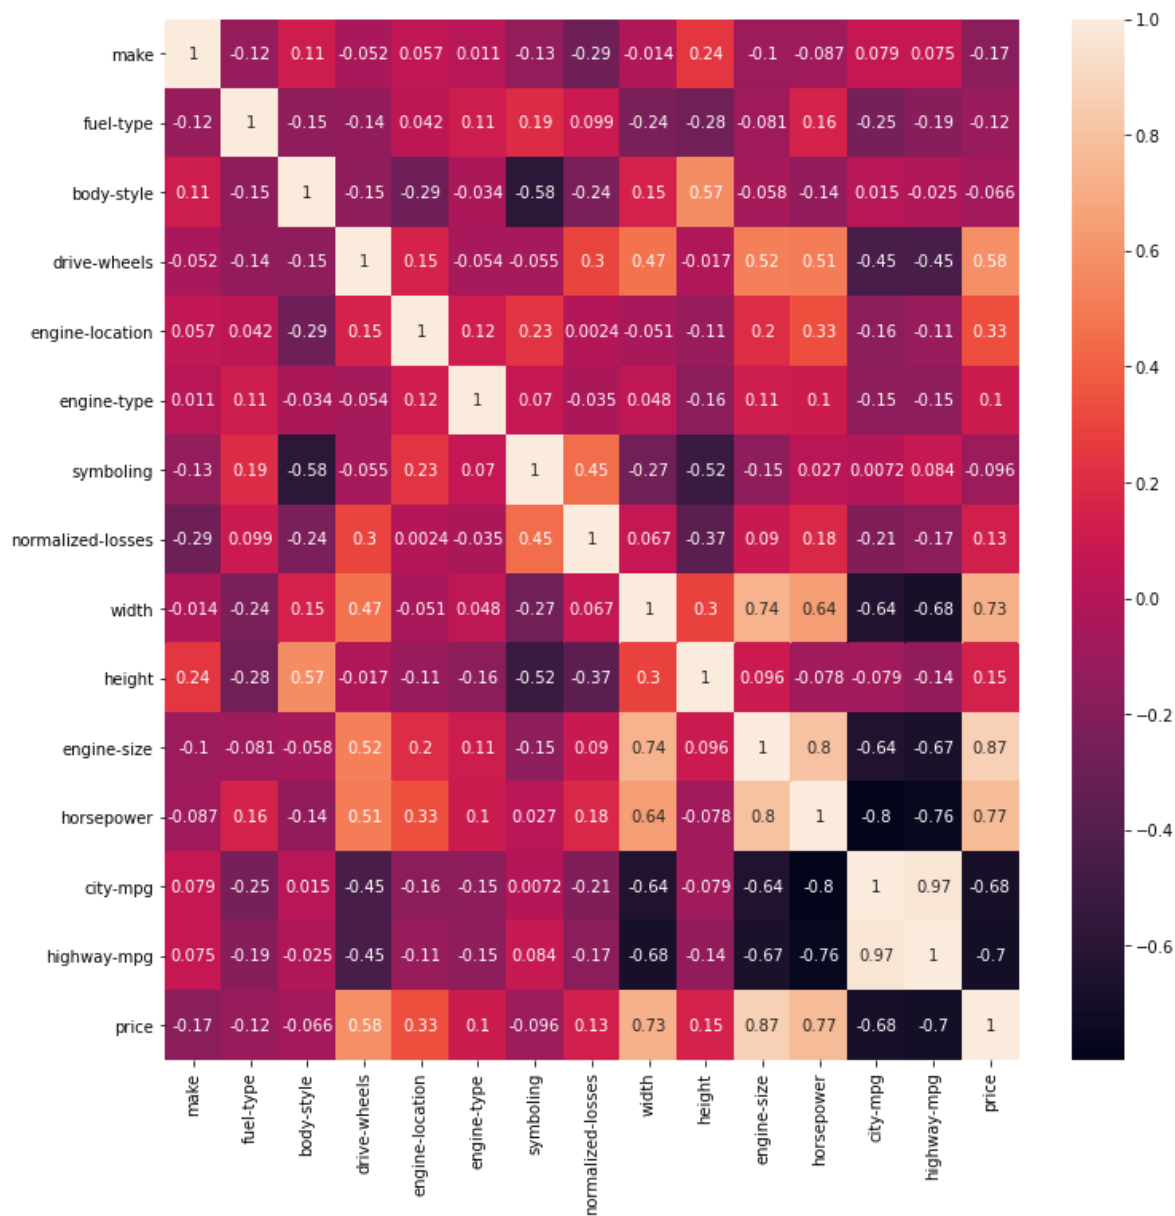
normalized-losses  
0.848205953606264

In [35]:

```
plt.figure(figsize=(12,12))
sns.heatmap(df.corr() , annot=True)
```

Out[35]:

&lt;AxesSubplot:&gt;



In [36]:

```
df.corr()["price"].sort_values()
```

Out[36]:

highway-mpg	-0.704846
city-mpg	-0.680412
make	-0.173792
fuel-type	-0.115791
symboling	-0.095905
body-style	-0.065831
engine-type	0.102758
normalized-losses	0.129973
height	0.147010
engine-location	0.333620
drive-wheels	0.584485
width	0.730503
horsepower	0.768921
engine-size	0.869638
price	1.000000

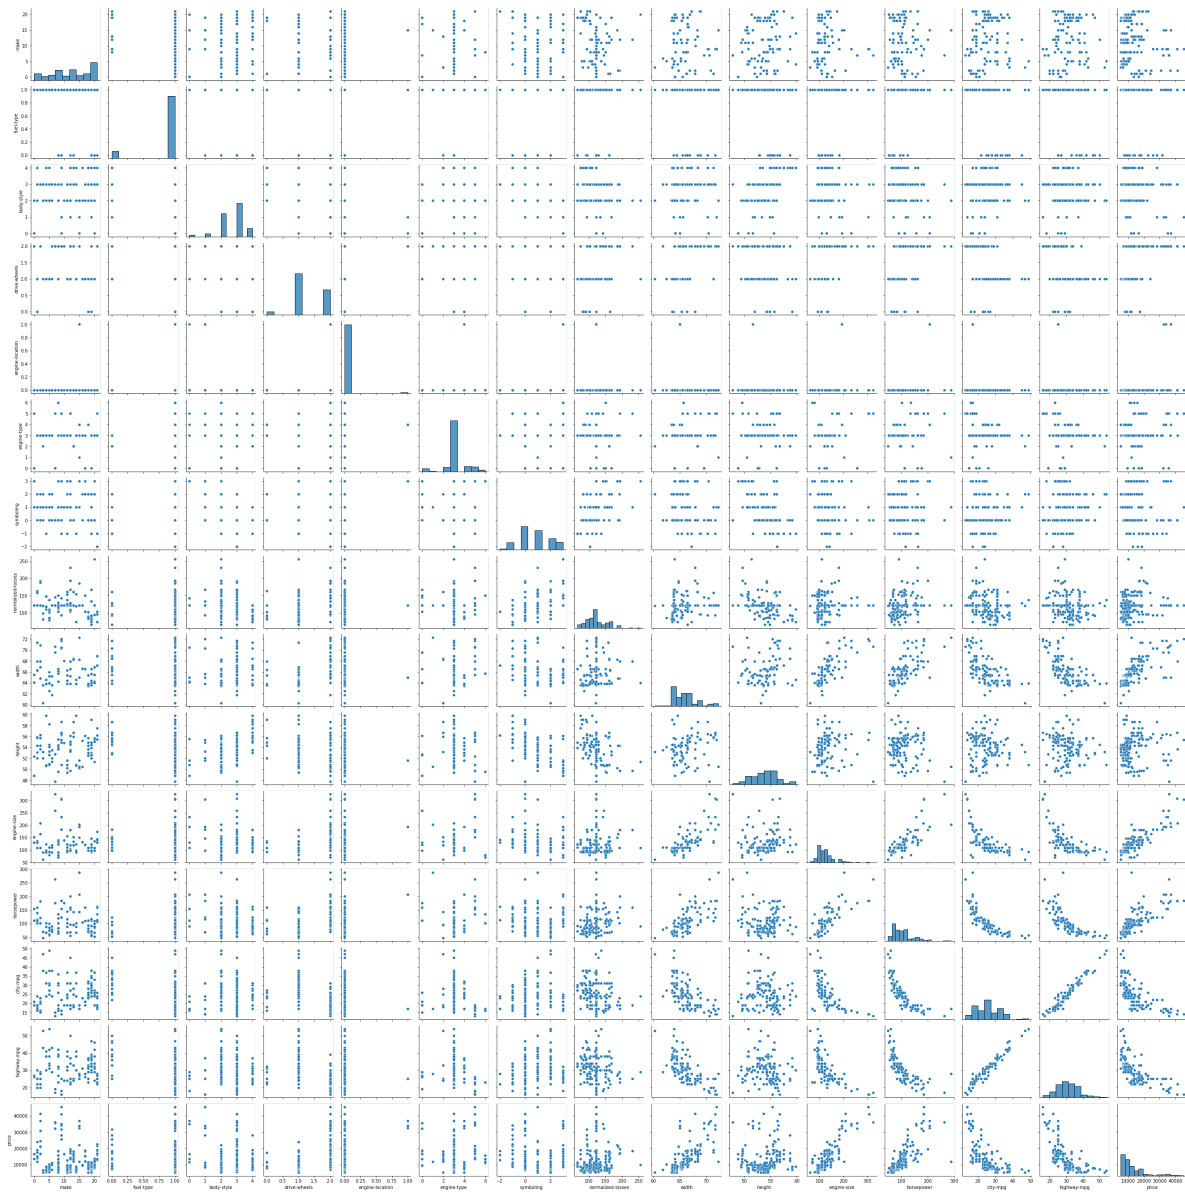
Name: price, dtype: float64

In [37]:

```
sns.pairplot(df)
```

Out[37]:

&lt;seaborn.axisgrid.PairGrid at 0x22476f2b520&gt;



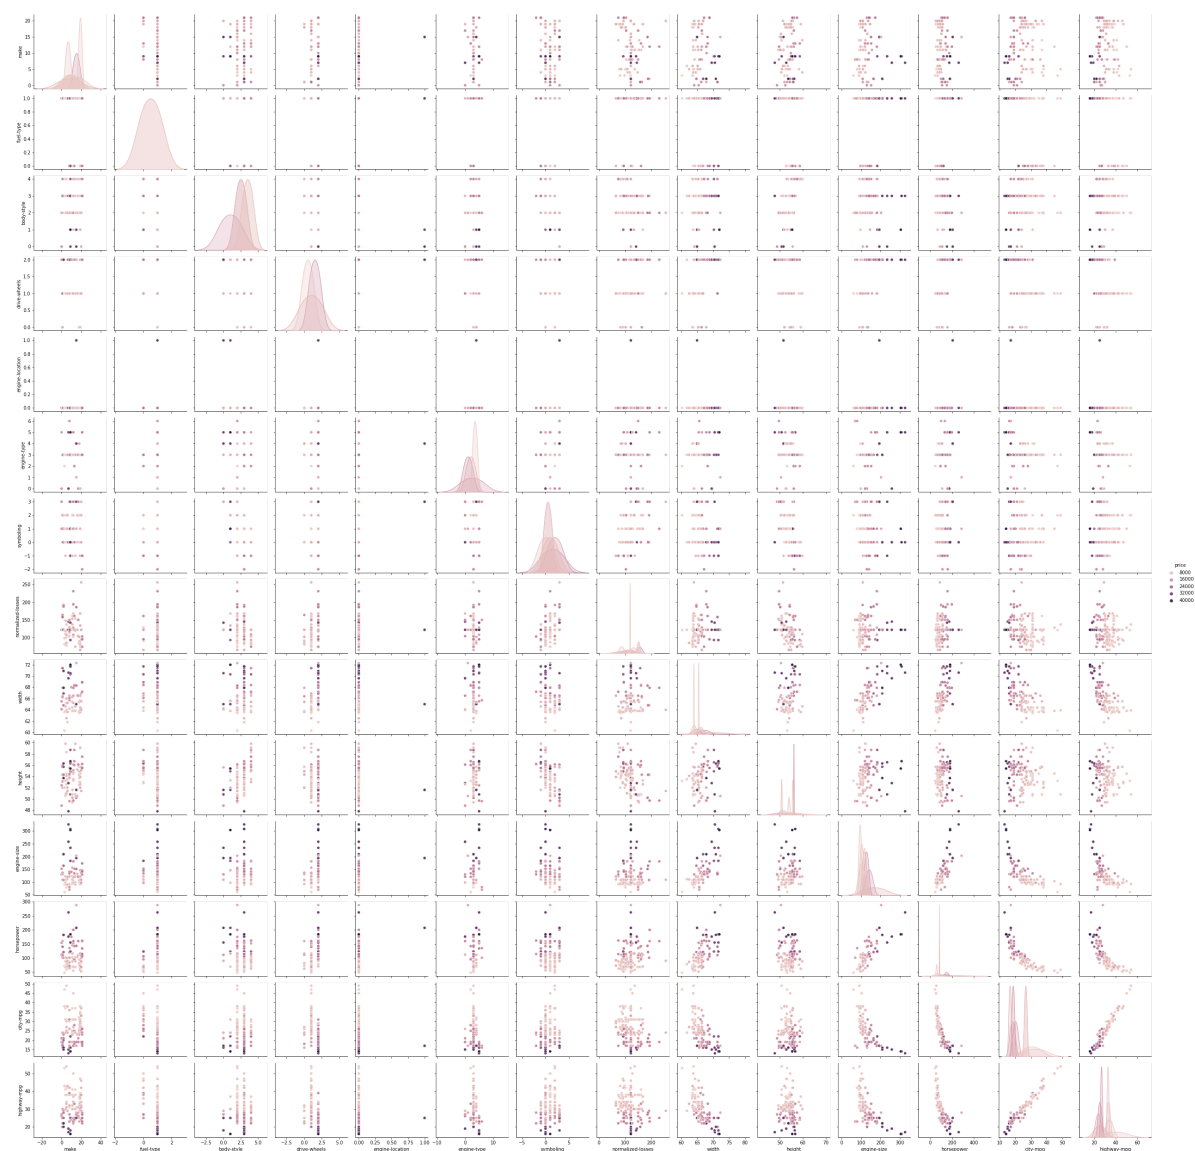


In [38]:

```
sns.pairplot(data=df , hue="price")
```

Out[38]:

```
<seaborn.axisgrid.PairGrid at 0x2247cc2f610>
```



## Model Creation

In [39]:

```
df.head()
```

Out[39]:

	make	fuel-type	body-style	drive-wheels	engine-location	engine-type	symboling	normalized-losses	width	height	engine-size
0	0	1	0	2	0	0	3	122.0	64.1	48.8	13
1	0	1	0	2	0	0	3	122.0	64.1	48.8	13
2	0	1	2	2	0	5	1	122.0	65.5	52.4	15
3	1	1	3	1	0	3	2	164.0	66.2	54.3	10
4	1	1	3	0	0	3	2	164.0	66.4	54.3	13

In [40]:

```
x = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

In [41]:

```
xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.2, random_state=1)
```

In [42]:

```
def mymodel(model):
    model.fit(xtrain, ytrain)
    ypred = model.predict(xtest)

    ac = r2_score(ytest, ypred)
    mae = mean_absolute_error(ytest, ypred)
    mse = mean_squared_error(ytest, ypred)
    rmse = np.sqrt(mse)

    train = model.score(xtrain, ytrain)
    test = r2_score(ytest, ypred)

    print(f"Training accuracy :- {train}")
    print(f"Testing accuracy :- {test}")
    print()
    print(f"R2_score :- \n{ac}\n\nMAE :- \n{mae}\n\nMSE :- \n{mse}\n\nRMSE :- \n{rmse}")
```

In [43]:

```
models = []

models.append(("Linreg", LinearRegression()))
models.append(("KNN", KNeighborsRegressor()))
models.append(("SVM-l", SVR(kernel="linear")))
models.append(("SVM-r", SVR(kernel="rbf")))

for name, model in models:
    print(name)
    mymodel(model)
    print("\n\n\n")
```

Linreg :-  
Training accuracy :- 0.8941185334320063  
Testing accuracy :- 0.7355527822266537

R2\_score :-  
0.7355527822266537

MAE :-  
3184.5155994121515

MSE :-  
21864056.153787263

RMSE :-  
4675.901640730615

KNN :-  
Training accuracy :- 0.9158165653923884  
Testing accuracy :- 0.6671791815444368

R2\_score :-  
0.6671791815444368

MAE :-  
3298.7692307692305

MSE :-  
27517071.743589748

RMSE :-  
5245.671715194323

SVM-l :-  
Training accuracy :- 0.8339318161071796  
Testing accuracy :- 0.6151716897711921

R2\_score :-  
0.6151716897711921

MAE :-  
3860.550037158177

MSE :-  
31816964.667864848

RMSE :-  
5640.652858301497

SVM-r :-  
Training accuracy :- -0.13936583233415778  
Testing accuracy :- -0.1791158963225512

R2\_score :-  
-0.1791158963225512

MAE :-  
5925.926834043423

MSE :-  
97487341.27774158


RMSE :-  
9873.56780894027

In [44]:

```
accuracy=[]  
  
for i in range(1,51):  
    knn = KNeighborsRegressor(n_neighbors=i)  
    knn.fit(xtrain , ytrain)  
    ypred = knn.predict(xtest)  
    accuracy.append(r2_score(ytest,ypred))
```

In [45]:

```
plt.plot(range(1,51),accuracy)  
plt.grid(True)  
plt.show()
```



In [46]:

```
n12 = []  
for i in range(100):  
    l2 = Ridge(alpha=i)  
    l2.fit(xtrain,ytrain)  
    n12.append((i,l2.score(xtest,ytest)))  
  
n12
```

Out[46]:

```
[(0, 0.7355527822266543),  
(1, 0.7016542422732281),  
(2, 0.6847205630664456),  
(3, 0.6751690951781149),  
(4, 0.6689336369274135),  
(5, 0.6644614007235592),  
(6, 0.6610439498857883),  
(7, 0.6583127977364198),  
(8, 0.6560570137239903),  
(9, 0.6541467733216799),  
(10, 0.6524975027598207),  
(11, 0.6510515309727207),  
(12, 0.6497680257188643),  
(13, 0.6486171525383218),  
(14, 0.6475765216649285),  
(15, 0.6466289392910423),  
(16, 0.6457609352149597),  
(17, 0.6449617700485919),  
(18, 0.6442227483265888),  
(19, 0.6435367323275265),  
(20, 0.6428977909174949),  
(21, 0.6423009412692836),  
(22, 0.6417419557549824),  
(23, 0.6412172154047594),  
(24, 0.6407235971861645),  
(25, 0.6402583862165862),  
(26, 0.6398192066102026),  
(27, 0.639403966428487),  
(28, 0.6390108134298362),  
(29, 0.6386380991777165),  
(30, 0.6382843496833958),  
(31, 0.6379482412052111),  
(32, 0.6376285801525738),  
(33, 0.6373242862842188),  
(34, 0.6370343785705443),  
(35, 0.636757963226026),  
(36, 0.6364942235212865),  
(37, 0.6362424110640186),  
(38, 0.6360018382995689),  
(39, 0.6357718720300747),  
(40, 0.6355519277888122),  
(41, 0.6353414649362952),  
(42, 0.6351399823684665),  
(43, 0.6349470147463903),  
(44, 0.6347621291722089),  
(45, 0.6345849222485911),  
(46, 0.6344150174690382),  
(47, 0.6342520628947241),  
(48, 0.6340957290803899),
```

```
(49, 0.633945707217447),  
(50, 0.6338017074671669),  
(51, 0.6336634574607254),  
(52, 0.6335307009461641),  
(53, 0.6334031965650795),  
(54, 0.6332807167441894),  
(55, 0.633163046688879),  
(56, 0.6330499834675272),  
(57, 0.6329413351768225),  
(58, 0.6328369201795169),  
(59, 0.6327365664071011),  
(60, 0.6326401107208057),  
(61, 0.6325473983250953),  
(62, 0.6324582822285081),  
(63, 0.6323726227472874),  
(64, 0.632290287047744),  
(65, 0.6322111487237436),  
(66, 0.6321350874061213),  
(67, 0.6320619884011252),  
(68, 0.6319917423553328),  
(69, 0.6319242449447335),  
(70, 0.6318593965858967),  
(71, 0.6317971021673566),  
(72, 0.6317372707995456),  
(73, 0.6316798155817325),  
(74, 0.6316246533846166),  
(75, 0.631571704647304),  
(76, 0.6315208931875617),  
(77, 0.6314721460242934),  
(78, 0.6314253932113333),  
(79, 0.6313805676816733),  
(80, 0.6313376051013773),  
(81, 0.6312964437324361),  
(82, 0.6312570243039552),  
(83, 0.6312192898910409),  
(84, 0.6311831858008601),  
(85, 0.6311486594653745),  
(86, 0.6311156603402739),  
(87, 0.6310841398096971),  
(88, 0.631054051096347),  
(89, 0.631025349176634),  
(90, 0.6309979907005325),  
(91, 0.6309719339158172),  
(92, 0.630947138596424),  
(93, 0.6309235659746519),  
(94, 0.6309011786769745),  
(95, 0.6308799406632299),  
(96, 0.6308598171689843),  
(97, 0.6308407746508724),  
(98, 0.6308227807347309),  
(99, 0.6308058041663689)]
```

In [47]:

```
n11 = []
for i in range(100):
    l1 = Lasso(alpha=i)
    l1.fit(xtrain,ytrain)
    n11.append((i,l1.score(xtest,ytest)))

n11
```

Out[47]:

```
[(0, 0.7355527822266535),
 (1, 0.7349043968321456),
 (2, 0.734224129247808),
 (3, 0.7335119482950561),
 (4, 0.7327678216112494),
 (5, 0.7319918069060299),
 (6, 0.7311838864554618),
 (7, 0.7303440009125781),
 (8, 0.7294722457664095),
 (9, 0.728568558705089),
 (10, 0.7276329142276274),
 (11, 0.7266654426694357),
 (12, 0.7256659278294073),
 (13, 0.7246344541902952),
 (14, 0.7235711776815625),
 (15, 0.7224758377877629),
 (16, 0.7213487164982898),
 (17, 0.7201895305351079),
 (18, 0.7189984282153095),
 (19, 0.7177754221166048),
 (20, 0.7165203766929641),
 (21, 0.71523345414556),
 (22, 0.7139146862983363),
 (23, 0.7125639314145946),
 (24, 0.7111812791847517),
 (25, 0.7097664480053986),
 (26, 0.7083200026249673),
 (27, 0.7068417949659476),
 (28, 0.7053387530630352),
 (29, 0.703825264895338),
 (30, 0.7022833223294923),
 (31, 0.7007156529267059),
 (32, 0.6991170851658983),
 (33, 0.6974860191261013),
 (34, 0.6958241390246818),
 (35, 0.6941295302583039),
 (36, 0.6924042214194559),
 (37, 0.690646107618992),
 (38, 0.6888562462795077),
 (39, 0.6870357825159619),
 (40, 0.6851823850675423),
 (41, 0.683298529611013),
 (42, 0.6813815856828743),
 (43, 0.6794343568301573),
 (44, 0.6774538635164875),
 (45, 0.6754432355182265),
 (46, 0.6734010488724398),
 (47, 0.6713249820637119),
```



```
(48, 0.6692193955744843),  
(49, 0.6670797874674177),  
(50, 0.6649084727382653),  
(51, 0.6627054026816726),  
(52, 0.6604705772976398),  
(53, 0.6582040808387541),  
(54, 0.6559060756859811),  
(55, 0.6535910737116393),  
(56, 0.6512443878924672),  
(57, 0.6488660296640226),  
(58, 0.6464561105279234),  
(59, 0.644014572224729),  
(60, 0.6415413778776521),  
(61, 0.6390365387265051),  
(62, 0.6365160593232179),  
(63, 0.6340051835687787),  
(64, 0.6320858674588254),  
(65, 0.6319941472925529),  
(66, 0.6319023132801276),  
(67, 0.6318125461834823),  
(68, 0.6317200169467454),  
(69, 0.6316273948684725),  
(70, 0.6315345696511616),  
(71, 0.6314438329521539),  
(72, 0.6313504617482983),  
(73, 0.6312569097328515),  
(74, 0.6311631558307014),  
(75, 0.6310694268918934),  
(76, 0.6309754456784029),  
(77, 0.6308836678692045),  
(78, 0.630789171403301),  
(79, 0.6306943923956252),  
(80, 0.6305993617142169),  
(81, 0.6305040879516217),  
(82, 0.6304086000844271),  
(83, 0.6303128605110933),  
(84, 0.6302168863052457),  
(85, 0.630120682703242),  
(86, 0.6300242273703448),  
(87, 0.6299275197257057),  
(88, 0.6298305598564977),  
(89, 0.6297333708063313),  
(90, 0.6296359391984844),  
(91, 0.6295382558267923),  
(92, 0.6294427720509725),  
(93, 0.6293567923373555),  
(94, 0.6292773020643387),  
(95, 0.6291975819920547),  
(96, 0.629117638518707),  
(97, 0.6290374703564884),  
(98, 0.6289570770217647),  
(99, 0.6288764657124297)]
```

In [48]:

```
models = []

models.append(("Linreg", LinearRegression()))
models.append(("KNN", KNeighborsRegressor(n_neighbors=1)))
models.append(("SVM-1", SVR(kernel="linear")))
models.append(("SVM-r", SVR(kernel="rbf")))

for name, model in models:
    print(name)
    mymodel(model)
    print("\n\n\n")
```

```
Linreg      :-
Training accuracy :- 0.8941185334320063
Testing accuracy  :- 0.7355527822266537
```

```
R2_score :-
0.7355527822266537
```

```
MAE :-
3184.5155994121515
```

```
MSE :-
21864056.153787263
```

```
RMSE :-
4675.901640730615
```

```
KNN      :-
Training accuracy :- 0.9928862905865443
Testing accuracy  :- 0.74027773895488
```

```
R2_score :-
0.74027773895488
```

```
MAE :-
2592.5384615384614
```

```
MSE :-
21473404.589743588
```

```
RMSE :-
4633.940503474725
```

```
SVM-1      :-
Training accuracy :- 0.8339318161071796
Testing accuracy  :- 0.6151716897711921
```

```
R2_score :-
0.6151716897711921
```

MAE :-  
3860.550037158177

MSE :-  
31816964.667864848

RMSE :-  
5640.652858301497

SVM-r :-  
Training accuracy :- -0.13936583233415778  
Testing accuracy :- -0.1791158963225512

R2\_score :-  
-0.1791158963225512

MAE :-  
5925.926834043423

MSE :-  
97487341.27774158

RMSE :-  
9873.56780894027

## Cross Validation Score

In [49]:

```
print("Name          Accuracy          STD")
for name, model in models:
    cvs = cross_val_score(model, x, y, cv=5,)
    print(f"{name} {cvs.mean()} {cvs.std()}")
```

Name	Accuracy	STD
Linreg	:- 0.5155626638216251	0.29417502732645967
KNN	:- 0.3998279418728215	0.3576296707751772
SVM-l	:- 0.5323432027166152	0.33526075240067404
SVM-r	:- -0.1881168806138271	0.13892274676036426

In [ ]:

